



HACKYEAH 2026 · AI OPEN TASK

You have free days. We tell you where and when.

Tell us your free days, budget and taste. Get ranked trips, every number sourced.

An auditable recommendation engine, not a chatbot.

Real screenshots of the live app · demo profile · Kraków

Proactive: a pick for your free dates



TRIPAI · now



You're free 7–11 Nov → Málaga

1,773 PLN all-in · flights 441 PLN · score 74. Tap to see why.

9:41



THE PROBLEM

Free days, no plan, and no answer you can check.



Ola, 29, Kraków

Free 7–11 Nov (long weekend)

Budget 1,800 PLN · Hates crowds

**Where should I go, and is this
even the right weekend?**

Travel agencies: opaque

A package and a price. No reasons.

Chat AI: unverifiable


Confident, unsourced. You redo the work.

89%



want AI help planning trips


12%



would let AI decide alone.
Trust is the gap.


THE SOLUTION

From "I'm free" to a trip you can check.




1 Travel DNA

14 swipes set your taste




2 Long weekends

2 days off -> 5 days away



3 Where + when

Ranked, priced all-in




4 Why this trip

Fit + sources, plainly

HOW IT WORKS

Numbers from sources. Ranking from a formula.

Every morning: one durable DBOS workflow, 07:00



GPT-6 Luna never produces a price or a score. It may only quote numbers from the evidence.

Built with pydantic-ai, FastAPI, DBOS, Supabase (row-level security), Next.js · first results in ~1.5 s, verdicts follow.

AI'S ROLE

AI only where it earns its place, measured.

Jev decides

Small TypeScript classifier: a decision + confidence, no prose.

A formula ranks

Sourced numbers in, same score out.

GPT-6 Luna: two jobs

① decides when Jev is unsure

② writes the "why" from evidence

Fit verdict vs 20 hand labels, at the low-cost tier			
ENGINE	EXACT MATCH	P50	\$ / VERDICT
Jev → GPT-6 Luna	16/20	3.0 s	\$0.00022
Jev alone	14/20	0.5 s	\$0.00011
GPT-6 Luna alone	14/20	6.5 s	\$0.00031
Rules only (no AI)	10/20	0 s	\$0

At this cost tier the cascade beats either model alone. GPT-6 Luna is OpenAI's cheapest model; stronger LLMs may score higher at more cost and latency. Our test: 20 hand-labelled cases, 3 Oct; n = 20, still labelling.


VERIFY AND STAY IN CONTROL

Every number has a receipt.

- ✓ **Source + timestamp** on every price and fact
- ✓ **Reproducible**: same inputs → same score
- ✓ **Budget first**: over-budget trips rank last, flagged
- ✓ **You decide**: nothing is booked without your tap
- ✓ **Personalisation off**: neutral weights, nothing learned


If an AI answer is hard for the user to check, that difficulty is a product smell.

Paraphrasing Hamel Husain (2026): the principle we built on



LEARNING + PROACTIVE

Learns visibly. Pings only when it matters.



Swipe → profile updates

Visible learning


Each swipe nudges your profile a little, visibly, and undo reverts it.

♥ Málaga: tags +0.05 · undo reverts

Proactive, not noisy

A durable scan of all your free windows every morning. Jev decides if a find is worth a push.

push only if $p(\text{worth it}) \geq 0.8$



p = 0.68 → inbox, no push

ChatGPT talks. TripAI computes, monitors, proves.

1 Where + when is a search problem

Hundreds of exact-date price checks across all your free windows: candidates, shortlist, cache, quota. First results in ~2 s. A chat agent samples a few.

2 Guarantees, not prose

Exact-date prices. The card equals the receipt: flight + hotel. Same inputs → same score.

3 Data generic tools don't join

Eurostat crowds, Polish bridge days and school holidays, climate normals for your exact dates.

4 Keeps working after the chat

A durable daily scan of your windows, and learning from your swipes and trip feedback.

5 And ChatGPT can call TripAI. Roadmap: an MCP server

find_free_windows

recommend_trips

get_trip_details

explain_trip

verify_ranking

watch_trip

update_travel_dna

book_links

BUSINESS MODEL · HONEST LIMITS

We earn on bookings and Premium, never on ranking.

HOW WE MAKE MONEY

Free

Recommendations, receipts, daily picks

Affiliate

Commission when you book through a partner

Premium ~19 PLN / month (hypothesis)

Price alerts, group trips; cost-splitting later

≈ **36 PLN**

per 1,800 PLN trip at 2%

≈ **0.004 PLN**

AI cost per ranking

AI cost ≈ 0.01% of the commission. Commission rate is an assumption.

HONEST LIMITS

Not a tour operator

→ partners book

Prices are quotes

→ re-checked before you pay

Demo calendar is a sample

→ Google connector built

Eval set: 20 cases

→ labelling more

Kraków + PL holidays

→ more origins next

BUILT, TESTED, LIVE

21

pull requests merged

365

backend tests

147

frontend tests

4

data sources

Live: flights, weather, AI · Seeded: crowds · Cached: hotels · Sample: calendar

Tested with 2 real users. Top finding: people don't read.

So we cut clutter, show value for money, added party size and honest price labels.

[Team name] · [Members]

AI use: built with Claude (Anthropic) coding agents, reviewed by the team; every PR AI-reviewed + smoke-tested. In the app: Jev (TypeSafe), GPT-6 Luna. Prices and scores never come from an LLM.

 **TripAI**

You have free days.
We tell you where and when.



<frontend-url>

Scan to try · no sign-up · demo profile

**Long weekend 7–11 Nov.
Where would you go?**